# STAT GU4206/5206 Midterm

*Your name and UNI goes here*

*June 8, 2017*

The STAT GU4206/5206 midterm is open notes, open book(s), and online resources are allowed. Students are **not** allowed to communicate with any other people during the midterm with the exception of the GU4206/5206 instructor. When you are finished with the midterm, please upload both the .pdf and .Rmd files on Canvas.

For the entire midterm we consider the **Auto** dataset taken from the *Introduction to Statistical Learning* package. Before starting the midterm, make sure the package **ISLR** is installed on your laptop.

```
#install.packages("ISLR")
library(ISLR)
```

## Problem 1: Basic Operations

### 1.i)

Display the first 3 rows of the **Auto** dataset:

```
head(Auto,3)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1  18         8          307        130   3504         12.0   70      1
## 2  15         8          350        165   3693         11.5   70      1
## 3  18         8          318        150   3436         11.0   70      1
##                       name
## 1 chevrolet chevelle malibu
## 2         buick skylark 320
## 3        plymouth satellite
```

### 1.ii)

How many rows are in this dataset?

```
# Solution
```

## Problem 2: Regression and the Bootstrap

### 2.i)

Consider extracting only the rows corresponding to 6 cylinder cars. Also consider running a linear regression on a car's acceleration versus its weight. The working filtering and linear regression code is displayed below.

```
Auto.6.cyl <- Auto[Auto$cylinders==6,]
beta.hats <- coef(lm(acceleration~weight,data=Auto.6.cyl))
beta.hats
```

```
## (Intercept)      weight
## 4.368123570 0.003711944
```

Identify the estimated slope and intercept of the above linear model.

**Solution:**

## 2.ii)

Now suppose as researchers, we want to infer upon the true slope relating a six cylinder car's acceleration versus its weight. Also suppose that we do not want to make strong assumptions on the errors of the linear regression model; hence we will perform a bootstrap procedure. Below is *almost complete* working code that runs a bootstrap procedure for the slope. Fill in the one missing line of **R** code and make sure to uncomment each line below. Run the bootstrap procedure after filling in the missing line.

```
# set.seed(0)
# B <- 1000
# n <- nrow(Auto.6.cyl)
# slopes.boot <- rep(NA,B)
# for (b in 1:B) {
#
#   regression.boot <- lm(acceleration~weight,data=Auto.6.cyl[sample.boot,])
#   slopes.boot[b] <- coef(regression.boot)[2]
#
# }
# slope.hat <- beta.hats[2]
# LL <- 2*slope.hat-quantile(slopes.boot,.975)
# UL <- 2*slope.hat-quantile(slopes.boot,.025)
# c(LL,UL)
```

## 2.iii)

At 5% significance, is a six cylinder car's acceleration statistically related to its weight? Support your answer using the computed bootstrap interval **c(LL,UL)**.

**Solution:**

## 2.iv)

Create a histogram of the bootstrapped slope estimates. Make sure to label the histogram appropriately and use 30 breaks for the bins.

```
# Solution
```

# Problem 3: Subseting

The original **Auto** datafrme consists of cars with 3,4,5,6 and 8 cylinders. Create a new dataframe named **Auto.new** that consists of only the cars with 4,6 and 8 cylinders. Check that the number of rows in this new dataframe is equal to 385.

```
# Solution
```

# Problem 4: Character Srings and Regular Expressions

## 4.i)

Look at the first few cases of the variable **name**.

```
head(Auto$name)
```

```
## [1] chevrolet chevelle malibu buick skylark 320
## [3] plymouth satellite        amc rebel sst
## [5] ford torino               ford galaxie 500
## 304 Levels: amc ambassador brougham ... vw rabbit custom
```

Notice that the first word in each string is the car's company, i.e., chevrolet, buick, toyota, etc... Append a new variable on the **Auto.new** dataframe named **company** that displays the company of each car. For example, if the **name** of the car is "chevrolet chevelle malibu", then the car's company should be "chevrolet". Show the first three observations in this new dataset.

Note: You might have to convert the factor variable back into a character variable.

```
# Solution
```

## 4.ii)

When the experimenter was recording the data, he entered a few typos for the car's company names, i.e., one case shows "toyouta" and another case shows "vokswagen". Fix these two typos in the **Auto.new** dataframe by using the **grep** function to find the location of the typos and then assigning new strings to these elements. After fixing the typos, create a table of the variable **company**.

```
# Solution
```

# Problem 5: The Apply Family

## 5.i)

Using the appropriate apply function, compute the maximum **horsepower** per **company**. Also sort this output.

```
# Solution
```

### 5.ii)

Using the appropriate apply function, compute the average value of quantitative variables **mpg**, **displacement**, **horsepower**, **weight** and **acceleration**. To save some time I provided the vector of character strings in the below code chunk.

```
variables <- c("mpg","displacement","horsepower","weight","acceleration")
# Solution
```

## Problem 6: R Base Graphics

Construct a base **R** plot that shows a car's acceleration ($Y$) versus a its weight ($X$) split by the number of cylinders in the car (4,6, and 8). Note that you should be using the **Auto.new** dataset. For full credit, create the scatter plot and split the data up by different colors to represent the number of cylinders. Also create a legend and label the plot appropriately.

For extra credit, plot regression lines for each subgroup, i.e., plot 3 least squares lines: one line for 4 cylinders, one line for 6 cylinders and one line for 8 cylinders.

```
# Solution
```